Segment CDP Warehouse Runbook

Purpose

Lightweight v0 implementation validating Segment CDP data from ingestion to analytics.

Python → CSV → Snowflake → dbt → Tableau

Data & Validation

Synthetic Segment events are generated using Python/Pandas and loaded into Snowflake.

Core fields:

`event_id | event_name | user_id | anonymous_id | sent_at | received_at | properties`

Validation focuses on:

- **Event arrival** — event counts by type.
- **Identity stitching** — relationship between `anonymous_id` and `user_id`.
- **Schema consistency** — event property names and casing.

dbt Models

Model	Purpose
`events_src`	Standardizes and deduplicates events
`users_dims`	User traits, first seen and last seen
`active_users_7d`	Identified users active in the last 7 days

Tableau MVP

- **Event Volume Over Time**
- **Product Views vs Purchases**

These demonstrate that the modeled data is available and usable for business analytics.

Operational Checks

Daily: event volumes, identity stitching, dbt model execution.

Weekly: property/schema consistency, unexpected event changes, purchases and active users.

Success Criteria

Events arrive consistently, anonymous and identified activity can be associated, the event structure remains consistent, and the modeled data supports business metrics.

Data arrives → identities are associated → events are modeled → metrics are visualized.